

# MapReduce Code

## Word Count example

# Designing the Mapper

- Define a mapper class that extends hadoop Mapper
- imported from Hadoop MapReduce library  
`org.apache.hadoop.mapreduce.Mapper;`
- format:
  - `Mapper<inKey, inValue, outKey, outValue>`
- has a method called map

```
public static class TokenizerMapper
    extends Mapper<LongWritable, Text, Text, IntWritable>{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context
        ) throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

# Designing the Mapper

- Input key and input value types depends on the used input format
- We will use TextInputFormat to generate inputs for the map
- key is the line offset (type is `LongWritable`) —> mapper input key = `LongWritable`
- value is the content of the line (type is `Text`) —> mapper input value = `Text`

```
public static class TokenizerMapper

    extends Mapper<LongWritable, Text, Text, IntWritable>{

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    @Override

    public void map(LongWritable key, Text value, Context context
        ) throws IOException, InterruptedException {

        StringTokenizer itr = new StringTokenizer(value.toString());

        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

# Designing the Mapper

- Output key is `Text` which is initialized with value equal to 1
- Output value is `IntWritable` which contains one of the words in the line
- context is used to write the intermediate key-value pair

```
public static class TokenizerMapper
    extends Mapper<LongWritable, Text, Text, IntWritable>{
    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    @Override
    public void map(LongWritable key, Text value, Context context
        ) throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

# Designing the Reducer

- Define the reducer class that extends hadoop Reducer
- must specify the type of the input/output key-value pairs
- override the reduce function
- Mapper output key = `Text` —> Reducer input key = `Text`
- Mapper output value = `IntWritable` —> Reducer input value = `IntWritable`

```
public static class IntSumReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {
    private IntWritable result = new IntWritable();

    @Override
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

# Designing the Reducer

- reduce method iterates over all values associated with one key in a for loop
- sum values to generate the word count

```
public static class IntSumReducer
{
    extends Reducer<Text, IntWritable, Text, IntWritable> {
        private IntWritable result = new IntWritable();

        @Override
        public void reduce(Text key, Iterable<IntWritable> values,
                           Context context) throws IOException, InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum);
            context.write(key, result);
        }
    }
}
```

# Designing the Reducer

- output (key, value)
  - key is equal to the input key (word) of type **Text**
  - value is the result of summing all values of type **IntWritable**
- types of the output keys and values must match the types configured for the job output
- use the context to write the final output

```
public static class IntSumReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {
    private IntWritable result = new IntWritable();

    @Override
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

# The Job Configuration/ job driver

- The configuration is important because it indicates how the job is submitted
- make an instance of the job, give it a name
- tell hadoop which jar to run, the one that contains our main class

```
public class WordCount {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");
        job.setJarByClass(WordCount.class);
        job.setMapperClass(TokenizerMapper.class);
        job.setCombinerClass(IntSumReducer.class);
        job.setReducerClass(IntSumReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

//Mapper
public static class TokenizerMapper
    extends Mapper<LongWritable, Text, Text, IntWritable>{
    ....
}

//Reducer
public static class IntSumReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {
    ....
}
```



# The Job Configuration

- configure the job mapper, combiner (optional), reducer classes
- combiner can be the reducer

```
public class WordCount {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");
        job.setJarByClass(WordCount.class);
        job.setMapperClass(TokenizerMapper.class);
        job.setCombinerClass(IntSumReducer.class);
        job.setReducerClass(IntSumReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }

    //Mapper
    public static class TokenizerMapper
        extends Mapper<LongWritable, Text, Text, IntWritable>{
        ....
    }

    //Reducer
    public static class IntSumReducer
        extends Reducer<Text, IntWritable, Text, IntWritable> {
        ....
    }
}
```

# The Job Configuration

- Set the type of the final output key-value pairs

```
public class WordCount {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");
        job.setJarByClass(WordCount.class);
        job.setMapperClass(TokenizerMapper.class);
        job.setCombinerClass(IntSumReducer.class);
        job.setReducerClass(IntSumReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }

    //Mapper
    public static class TokenizerMapper
        extends Mapper<LongWritable, Text, Text, IntWritable>{
        ....
    }

    //Reducer
    public static class IntSumReducer
        extends Reducer<Text, IntWritable, Text, IntWritable> {
        ....
    }
}
```

# The Job Configuration

- Input / output configuration
  - input format class
    - here we use TextInputFormat class
  - output format class
    - TextOutputFormat class
- path pointing the the input data
  - read from the command line
- path to where final output should be written
  - read from the command line

```
public class WordCount {

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "word count");
        job.setJarByClass(WordCount.class);
        job.setMapperClass(TokenizerMapper.class);
        job.setCombinerClass(IntSumReducer.class);
        job.setReducerClass(IntSumReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        job.setInputFormatClass(TextInputFormat.class);
        job.setOutputFormatClass(TextOutputFormat.class);
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new Path(args[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

//Mapper
public static class TokenizerMapper
    extends Mapper<LongWritable, Text, Text, IntWritable>{
    ....
}

//Reducer
public static class IntSumReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {
    ....
}
```

# The Job Configuration

- Submit the job and wait until it finishes
- To launch the job
  - create a jar that contains the code

```
> hadoop jar "jar-file-name" "main  
class" "path-to-input" "output-path"
```

```
public class WordCount {  
  
    public static void main(String[] args) throws Exception {  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "word count");  
        job.setJarByClass(WordCount.class);  
        job.setMapperClass(TokenizerMapper.class);  
        job.setCombinerClass(IntSumReducer.class);  
        job.setReducerClass(IntSumReducer.class);  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        job.setInputFormatClass(TextInputFormat.class);  
        job.setOutputFormatClass(TextOutputFormat.class);  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}  
//Mapper  
public static class TokenizerMapper  
    extends Mapper<LongWritable, Text, Text, IntWritable>{  
    ....  
}  
//Reducer  
public static class IntSumReducer  
    extends Reducer<Text, IntWritable, Text, IntWritable> {  
    ....  
}
```

# Writables

- Hadoop writers for serializing/deserializing input/output data
- Hadoop has implementation for basic Java types
  - int —> IntWritable
  - long —> LongWritable
  - float —> FloatWritable
  - double —> DoubleWritable
  - Boolean —> BooleanWritable
  - String —> Text

# Summary of the Hadoop MapReduce job

- InputFormat
  - get the input data process into key-value pairs for the mapper
- Mapper
  - process each key-value pair & produce intermediate key-value pairs
- Combiner
  - an instance of the reducer that run as part of the map task to reduce intermediate output
- Partitioner
  - which key go to which reducer
- Reducer
  - take a key and all its values, process, and produce final output
- OutputFormat
  - write final output from reducers to HDFS

# References

- Google Research Paper
  - MapReduce: Simplified Data Processing on Large Clusters
- Apache Hadoop:
  - <http://hadoop.apache.org/>
- Book
  - Hadoop The Definitive Guide